

Bulk inserts with PostgreSQL

Four methods for efficient data loading

Ryan Booz, Developer Advocate

December 1, 2021

Agenda

- 01 Death by a thousand INSERT's
- 02 4 methods for bulk loading data
- 03 Demo
- 04 SDK Frameworks



01

Death by a thousand

INSERTs

**How often do you load
a lot of data into
PostgreSQL?**



01011

10011

00001

11001

10110

{JSON}

C, S, V




```
# "reader" = 1,000 rows
for row in reader:
    cur.execute(
        "INSERT INTO test_insert VALUES (%s, %s, %s)",
        row
    )
```



SLOW!






INSERTs are slow because...


- ...every statement incurs overhead
 - Network
 - Parsing
 - Planning
 - Locks
 - Execution
 - Response
- Even local application --> DB has overhead















02

Four methods for bulk loading data

Multi-valued INSERT



Multi-valued INSERT pseudocode

```
INSERT INTO table VALUES (1,2,3)
, (4,5,6)
, (7,8,9)
, (10,11,12)
, ... ;
```




```
# "reader" = 1,000 rows
sql = "INSERT INTO test_insert VALUES
next(reader) # Skip the header row.
for row in reader:
    batch_count += 1
    sql += "('{}', {}, {}),".format(*row)

    if batch_count == 500:
        cur.execute(sql[:-1])
        batch_count = 0
        sql = "INSERT INTO test_insert VALUES "
```



Multi-valued INSERT

- Little extra programming effort
- Supported regardless of the driver or language (even dynamic pl/pgsql)
- Moderately faster
- Generally requires batching subsets of rows

ARRAY values INSERT



ARRAY values INSERT pseudocode

```
INSERT INTO test_insert(time,tempc,cpu)
  SELECT * FROM unnest(
    $1::timestampz[],
    $2::integer[],
    $3::float8[]
  ) a(a,b,c)
ON CONFLICT DO NOTHING;
```



```
# converting column data to list
date = data['time'].tolist()
tempc = data['tempc'].tolist()
cpu = data['cpu'].tolist()

i=0
batch_size=2000
batch_end=batch_size
total_length = len(date)
while batch_end < total_length:
    cur.execute(
        "INSERT INTO test_insert SELECT * FROM
        unnest(%s::timestampz[],%s::int[],%s::double precision[])
        a(t,v,s)
        ON CONFLICT DO nothing;",
        (date[i:batch_end],tempc[i:batch_end],cpu[i:batch_end]))
    i=batch_end
    batch_end+=batch_size
    conn.commit()
```



ARRAY values INSERT

- Can Faster than multi-valued INSERT in some cases
- Avoids 65535 parameter limit 😊
- YMMV with language support for PostgreSQL arrays
- Caution: May not handle custom types correctly

COPY



COPY*

- Preferred, optimized tool for PostgreSQL bulk load
- Reads from files or STDIN
- Paths are local to PostgreSQL server
- Can also pull data out to a file
- Not in the SQL standard

*Not **psql \copy** (but closely related)



COPY Limitations

- Single transaction
- Single threaded
- No progress feedback prior to PG14
 - **pg_stat_progress_copy** view
 - In **psql**, run every second: **SELECT * FROM pg_stat_progress_copy \watch 1**
- No failure tolerance - stops on first error
 - Failed import takes space and leaves rows inaccessible 😞
- Minimal format configuration

**COPY does one job
extremely well: importing/
exporting data fast!**

Unlogged Tables



UNLOGGED Tables - Caution!

- Data inserted into UNLOGGED tables is not written to the Write-ahead Log (WAL)
- Eliminates some of that INSERT overhead
- Data is not crash safe
- Not accessible on replication servers (requires WAL)
- Available with CREATE and ALTER table



Why use UNLOGGED tables?

- You data process can accept the risk of loss for increased INSERT
- Staging tables for ETL processes
- Intermittent, repeatable work (easy to redo)

03

Demo

04

SDK Frameworks



What to look for in language SDKs

- Specific COPY support (including BINARY COPY format)
- Multi-valued and batching functions
- Is auto-commit enabled by default?
- Avoid parameterized query formatter
 - Use the ARRAY trick if the SDK only uses parameterized queries

**What questions do
you have for me?**

Thanks!

Me: **Ryan Booz**

Twitter: **@ryanbooz**

Slack: **slack.timescale.com**

