

Point-in-time query tuning and observability with `pg_stat_statements`

Ryan Booz (@ryanbooz)

Director, Developer Advocacy

April 2022



Timescale

Agenda

- 01** pg_stat_statements primer
- 02** How to track historical data
- 03** Demo
- 04** Other alternatives



pg_stat_statements primer



What is pg_stat_statements?

- An extension included with PostgreSQL 8.4+
- It is part of the ***contrib*** module but **not enabled** by default
 - Must be loaded via 'shared_preload_libraries' in postgresql.conf
- Tracks aggregated statistics of all queries in the cluster
- Installing the extension in the database creates the necessary views to query the data



How does it store aggregates?

- Every dbid, userid, and queryid
- Stats are grouped based on query structure and final ID as determined by an internal hash calculation



How does it identify queries?

```
SELECT id, name FROM table1 WHERE id = 1000;
```



```
SELECT id, name FROM table1 WHERE id = $1;
```

```
SELECT id, name FROM table1 WHERE id IN  
(1000,2000,3000);
```



```
SELECT id, name FROM table1 WHERE id IN  
($1,$2,$3);
```



pg_stat_statement statistics

- Execution Time (total/min/max/mean/stddev)
- Planning Time (total/min/max/mean/stddev)
- Calls (total)
- Rows (total)
- Buffers (shared/local/temp)
 - read/hit/dirtied/written
 - read/write time
- WAL

31 Columns of data
(as of PG14)



Name	Value
userid	16422
dbid	16434
queryid	-6155333619461995114
query	SELECT id, name FROM...
plans	0
total_plan_time	0.0
min_plan_time	0.0
max_plan_time	0.0
mean_plan_time	0.0
stddev_plan_time	0.0
calls	151
total_exec_time	8.489053
min_exec_time	0.013751
max_exec_time	1.356096
mean_exec_time	0.056218894039735096
stddev_exec_time	0.11851139585068957
rows	151
shared_blks_hit	450
shared_blks_read	3



**All statistics are cumulative
from the last restart***

***or reset by a superuser**

Caveats

- **PostgreSQL 13**
 - modified column names to include planning statistics
- **PostgreSQL 14**
 - Must set “compute_query_id”=true in postgresql.conf
 - Includes informational view for allocation and “last reset” information



02

Track historical data



(1) Create snapshot table(s)

- Store all or partial statistics
- Create additional columns from joined tables (ie. `pg_roles` or `pg_database`)
- Separate database can be helpful to easily filter out monitoring queries
- Partitioning highly recommended
 - TimescaleDB Hypertable provides automatic time partitioning and 90%+ compression of `pg_stat_statements` data



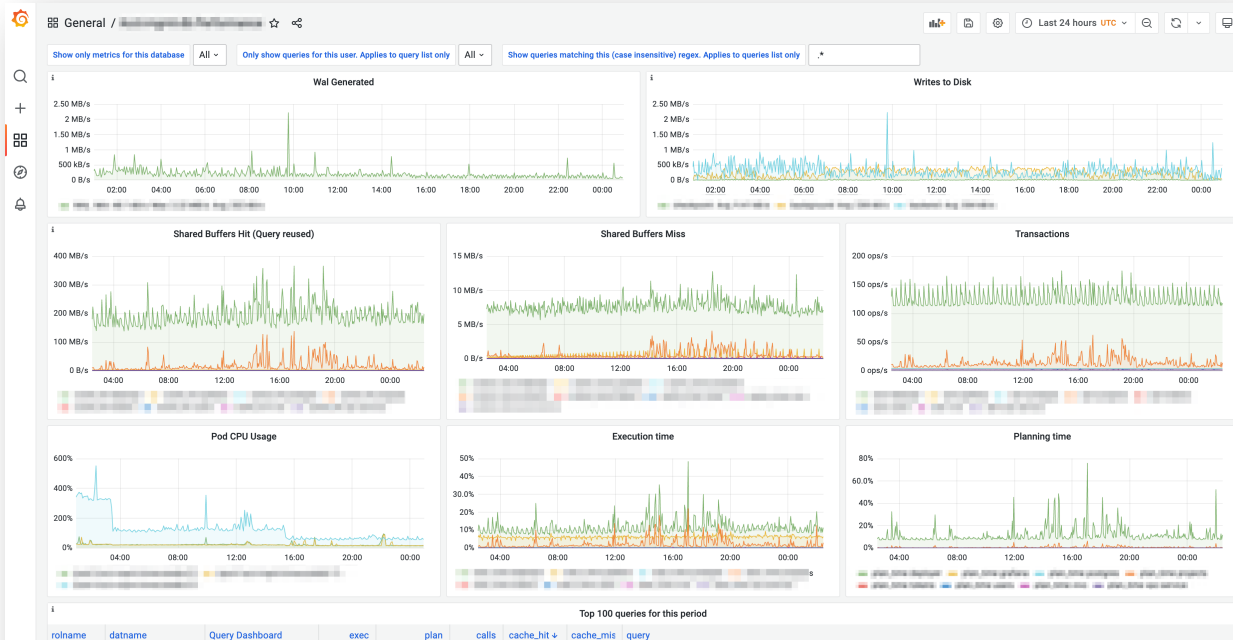
(2) Automate procedure to store statistics

- Create a stored procedure to query the view and store relevant data
- Potentially customized per database
- Automate the request through multiple tools or extensions
 - TimescaleDB User Defined Actions
 - pg_cron
 - pg_timetable



(3) Create VIEW to query snapshot differentials

- Monitor recent data through dashboards
- Show query specific values for a specific interval of time



03

Demo

04

Alternative solutions

Alternative Solutions

- **pg_stat_monitor**
 - Percona open source extension
 - Captures actual parameters
 - Query plans
 - Table access statistics
- **pganalyze**
 - Paid service to collect and analyze `pg_stat_statement` data
 - Additional log analysis tools as well



Thank you!

**What questions do
you have?**